

# STACMR

STACMR is a set of Matlab functions that conduct CMR analyses of state-trace data. CMR stands for “coupled monotonic regression”. This is a quick guide to its use.

## Java

The current implementation of STACMR calls code written in java. While this speeds up the analysis to a very great extent, it does require (at this point) some extra housekeeping. This is itemized below:

1. Java is usually pre-installed in Matlab. To check, type:

```
>> version -java
```

Any version on or after version 1.7 is good. If java is not installed, you will have to do this yourself from <http://www.oracle.com/technetwork/java/javase/downloads/jre8-downloads-2133155.html>.

2. The appropriate CMR java runtime library also needs to be acquired and stored in a convenient directory. The current version is called `fxMR-0.3.25.jar`.
3. Before running any of the STACMR programs, the CMR java runtime library needs to be made available to Matlab. This requires the command:

```
javaclasspath(path);
```

where path is the full pathname of the library. An example (omitting stuff in the middle) is:

```
javaclasspath('C:\Users\LAPTOP\Documents\...\fxMR-0.3.25.jar');
```

## Input data structures

The basic idea is that we have a set of data (called a *data structure*) that is organised as a set of observations on  $d > 1$  *dependent variables* across levels of two or more *independent variables*. From the point of view of *state-trace analysis* (STA), the number of different independent variables is not important – rather, it is the total number of conditions (combinations of levels). For example, if there are two independent variables, one with 2 levels, the other with 3 levels, then there will be a total of  $2 \times 3 = 6$  conditions. These conditions will be numbered from 1 to 6.

For this reason, STACMR distinguishes only between conditions defined between-subjects and conditions defined within-subjects. If there are  $b$  between-subjects conditions and  $w$  within-subjects conditions then there is a total of  $bw$  conditions ordered by the set of between-subjects conditions. That is, the first  $w$  conditions correspond to the set of within-subjects conditions under the first between-subjects condition, the next  $w$  conditions

correspond to the set of within-subjects conditions under the second between-subjects condition, and so on.

At present, STACMR accepts two kinds of data structure:

1. **Cell array format.** In this format, the data are organised in a  $b \times d$  cell array where  $b$  is the number of between-subjects groups and  $d$  is the number of dependent variables. Each component of this cell array is itself an  $n \times w$  matrix of observations where  $n$  is the number of subjects (which may vary across groups and dependent variables) and  $w$  is the number of within-subjects conditions (fixed across groups and dependent variables). The dependent variable may be either within-subjects or between-subjects – it doesn't matter because the correlation between dependent variables is assumed to be zero (although this might be up for grabs in future implementations).
2. **General format.** This structure is useful if the data are already in some kind of fixed column format. It is organised as a matrix in which each row corresponds to an observation and each column is defined as follows:  
column 1 = subject number (for identification only, not used directly)  
column 2 = between-subjects condition or group (if none, then column 2 = 1)  
column 3 = dependent variable (1 or 2)  
columns 4 to end = values for each within-subjects condition

While STACMR accepts data in general format it always converts it to cell array format using the function,

```
>> y = gen2cell (data);
```

### ***Partial order***

As well as data, the various STACMR functions often make use of an optional partial order. A partial order on a vector,  $x$ , is a set of pairs,  $(i, j)$ , such that  $x_i < x_j$ .

A partial order is represented in STACMR in two (equivalent) ways:

1. As a cell array containing the set of  $(i, j)$  pairs, e.g.,  $\{ [1 \ 2] \ [2 \ 3] \ [1 \ 4] \}$ . There is a shorthand for a linear order such as  $\{ [1 \ 2] \ [2 \ 3] \}$  which can be written as  $\{ [1 \ 2 \ 3] \}$  or, even more simply, as  $\{ 1:3 \}$ .
2. As an adjacency matrix in which entry  $(i, j) = 1$  if  $(i, j)$  is an element of a partial order, otherwise  $(i, j) = 0$ .

The function, `cell2adj`, converts a partial order in cell array form into its corresponding adjacency matrix form. For example,

```
>> E = { [1 2] [2 3] [1 4] };  
>> A = cell2adj (1:4, E)  
A =
```

0	1	0	1
0	0	1	0
0	0	0	0
0	0	0	0

In the above call to `cell2adj`, the vector, `1 : 4`, specifies the set of nodes or points that the partial order applies to. It is almost invariably the sequence of numbers, 1 to  $n$ , where, in this case,  $n = 4$ . The function `adj2cell` converts an adjacency matrix into its corresponding cell array.

### ***The principal functions***

The operation of the principal functions will be illustrated with respect to two data sets.

The first, called `delay`, is a data set discussed by Dunn, J. C., Newell, B. R., & Kalish, M. L. (2012). The effect of feedback delay and feedback type on perceptual category learning. *Journal of Experimental Psychology: Learning, Memory and Cognition*, 38(4), 840-859. It is in general format and is contained in the file, `delay.dat`. That is,

```
>> delay = load('delay.dat'); % read the data file
>> y = gen2cell(delay); % convert to cell array format
```

### ***staSTATS.m***

This function computes summary statistics of a data structure in cell array format.

Example call:

```
>> delaystats = staSTATS (y, shrink);
```

Here, `y` is the `delay` data in cell array format and `shrink` is an optional parameter denoting how much shrinkage to apply to the estimated covariance matrix. Generally, the covariance matrix needs to be shrunk during the bootstrap cycle to avoid ill-conditioning.

If `shrink = 0` then no shrinkage is applied. If `shrink = 1` then maximum shrinkage is applied. This means that the covariance matrix is diagonalized with all off-diagonal entries set to zero. If `shrink < 0` (the default) then an optimal shrinkage value is estimated for each within-subjects block and applied according to an algorithm developed by Ledoit, O. & Wolf, M. (2004). Honey, I shrunk the sample covariance matrix, *The Journal of Portfolio Management*, 30(4), 110-119.

`staSTATS` returns, in `delaystats`, a cell array of length equal to the number of dependent variables in `y`. Each component of `delaystats` is a structured array. For the dependent variable, `ivar`:

- `delaystats{ivar}.means` = vector of means across all conditions
- `delaystats{ivar}.cov` = the covariance matrix (for information only)
- `delaystats{ivar}.regcov` = the adjusted covariance matrix following application of shrinkage
- `delaystats{ivar}.n` = matrix of number of observations (subjects) in each within-subjects block
- `delaystats{ivar}.lm` = matrix of Loftus-Masson within-subjects standard errors (used by `staPLOT` below)
- `delaystats{ivar}.weights` = matrix of weights defined by:  
`delaystats{ivar}.n.* delaystats{ivar}.regcov^-1`
- `shrinkage` = a vector of length  $b$  (where  $b$  is the number of levels of the between-subject independent variable) containing the specified or estimated shrinkage values

Thus, `delaystats` has two elements. If you look at the means, they look like this:

```
>> disp(delaystats{1}.means)
0.3676  0.4676  0.5757  0.6118  0.3445  0.4434  0.5081  0.5169

>> disp(delaystats{2}.means)
0.3308  0.4550  0.5346  0.5492  0.2836  0.3031  0.3180  0.3098
```

### ***staMR.m***

This function conducts monotonic regression on a data structure according to a given partial order. We say it fits the *partial order model* to the data (i.e., the set of dependent variables).

Example call:

```
>> [x, f, shrinkage] = staMR (data, E, shrink);
```

#### ***Input:***

Here, `data` is either a data structure (in cell array or general format) or structured output from `staSTATS`; `E` is a partial order (required) in either cell array or adjacency matrix format; `shrink` is an optional shrinkage parameter (defined previously). If `data` is a cell array of structured output from `staSTATS`, then the shrinkage specified by this output is used whether the argument `shrink` is specified or not.

#### ***Output:***

`x` is a  $d$ -element cell array of that contains the best-fitting values for each dependent variable; `f` is the value of the least squares fit; `shrinkage` is a  $b \times d$  matrix of shrinkage values (where  $b$  is the number of levels of the between-subject independent variable).

Try this with `delay`. To do so, we have to specify a partial order. Use the following:

```

>> E = {1:4, 5:8, [5 1], [6 2], [7 3], [8 4]};

>> [x2, f2, s2] = staMR (delay, E);

>> disp([x2{:}]);
    0.3676    0.3308
    0.4676    0.4550
    0.5757    0.5346
    0.6118    0.5492
    0.3445    0.2830
    0.4434    0.3034
    0.5081    0.3149
    0.5169    0.3149

>> disp(f2);
    0.1721

>> disp(s2);
    0.0520    0.0491
    0.0314    0.2650

```

### ***staCMR.m***

This is the main function that conducts the CMR (state-trace) analysis. It takes a data structure or a cell array of structured output from `staSTATS` and an optional partial order and returns the best fitting values (to the data means) and the least squares fit. We say it fits the *conjoint partial order and monotonic model* to the data.

Example call:

```

>> [x, f, s] = staCMR (data, E, shrink);

```

On the input side, `data` is the data structure, `model` is an optional model (explained above), `E` is an optional partial order, `shrink` is an optional shrinkage parameter (defined previously).

On the output side, `x` is a cell array of the best-fitting values, `f` is the value of the least squares fit, and `s` is a structured array of fit statistics. The component, `s.shrinkage`, is a  $b \times d$  matrix of shrinkage values.

Now try this with `delay`:

```

>> [x1, f1, s1] = staCMR (delay, E);

>> disp([x1{:}]);
    0.3759    0.3150
    0.4850    0.4358

```

0.5898	0.5167
0.6265	0.5318
0.3353	0.2856
0.4186	0.3150
0.4806	0.3227
0.4850	0.3227

```
>> disp(f1);
1.7493
```

```
>> disp(s1.shrinkage);
0.0520    0.0491
0.0314    0.2650
```

### ***staCMRFIT.m***

This function estimates the empirical distribution (and hence  $p$ -value) of the fit of the conjoint monotonic and partial order model against the fit of the partial order model.

#### Example call:

```
>> [p, datafit, fits] = staCMRFIT (nsample, data, model, E,
shrink, proc);
```

#### *Input:*

- `nsample` is the number of Monte-Carlo samples (about 10,000 should be good)
- `data` is a data structure (which can be in general format, cell array format, or structured output from `staSTATS`). If data is non-summary data in general or cell array format then the bootstrap resampling is non-parametric, otherwise it is parametric (assumes a normal distribution).
- `model` is an optional model (default = STA)
- `E` is an optional partial order (default = none)
- `shrink` is an optional shrinkage parameter (default = -1)
- `proc` is an optional specification of maximum number of processors to be used (default = maximum available)

#### *Output:*


- `p` is the estimated  $p$ -value
- `datafit` is the observed fit value (see below)
- `fits` is a vector of empirical fit values (see below)

#### **Important note:**

staCMRFIT operates on the difference in fit values between a partial order model (calculated using staMR.m, e.g., f2 above) and the combined partial order + monotonic model (calculated using staCMR.m, e.g., f1 above).

Example output (note that p and fits will be approximate only):

```
>> [p, datafit, fits] = staCMRFIT (10000, delay, [], E);  
>> disp([p, datafit])  
    0.4538    1.5772  
>> histogram (fits, 100)
```



### ***staPLOT.m***

This function generates a state-trace plot of a data structure.

#### Example call:

```
>> staPLOT (data, optional argument pairs) ;
```

Here, data is a data structure or a cell array of structured output from staSTATS. The optional argument pairs are of the form; '*keyword*', *value*. The following are the set of keywords (valid alternatives in parentheses);

- predictions (prediction, pred, p) is a *d*-element cell array that defines a set of fitted values (output from staMR or staCMR)
- vars (v, dv) is a 2-vector that defines two dependent variables for plotting (default = [1, 2] )
- groups (group, g) is a cell array that defines a structure on the conditions. For example if there are 2 between-subjects conditions and 3 within subjects conditions then one group structure might be {1:3, 4:6}. At present, only one factor can be identified in this way.
- labels (label, lab, l) is a cell array that defines the labels of the levels defined by groups
- axislabels (axislabel, axislab, axis) is a cell array that defines the labels of the x- and y-axes
- axislimits (axislimit, axislim, limits, limit, lim) is a cell array that defines the upper and lower limits of the axes
- axisticks (axistick, ticks, tick, t) is a cell array that defines tick marks for x- and y-axes (if specified then axislimits is redundant)
- location (loc) is a Matlab value for the location of the legend (default = 'northwest')
- color (col, c) is a 3-element RGB vector specifying the color of the plot of predicted values (default is black)
- line is a flag to indicate if the predicted values are connected by a line. If set a line is drawn (default) otherwise no line is drawn.

#### Example output:

To see how this works, we use delay and the results of the earlier call to staCMR.

```
>> staPLOT(delay, 'pred', x1, 'groups', {1:4 5:8}, 'labels',  
{ 'Delay' 'No Delay'}, 'axislabels', {'RB' 'II'}, 'axisticks',  
{.2:.1:.7, .2:.1:.7});
```

This should produce the following graph:



